
P-439

CRUDE OIL PRICE PREDICTION

GROUP MEMBERS:

- GULAM MOHIUDDIN KHAN
 - MOHDGOUS SHAKEEL KURNE
 - NAMITHA KUAMRI
 - SHREEYA NITTURKAR
 - SOORA SHIRISHA
 - SURENDRA KUMAR ABALUR
-

OBJECTIVE

To forecast oil prices, which are often influenced by external factors rather than real-time data, making predictions challenging. Our model aims to identify patterns in oil price movements to assist customers and businesses in making informed decisions, ultimately mitigating the economic impact of fluctuating oil prices

DATA SET DESCRIPTION

The dataset named Crude oil Prices Daily.xlsx contains data collected to analyze and predict crude oil prices. The dataset includes various market factors and external influences that affect the fluctuations in oil prices, providing insights into the complex dynamics of global oil markets.

- **Number of Instances (Rows): 8223**
- **Number of Variables (Columns): 2**

	Date	Closing Value
0	1986-01-02	25.56
1	1986-01-03	26.00
2	1986-01-06	26.53
3	1986-01-07	25.85
4	1986-01-08	25.87
...
8218	2018-07-03	74.19
8219	2018-07-04	NaN
8220	2018-07-05	73.05
8221	2018-07-06	73.78

- **Data Source:** The dataset, named Crude Oil Prices Daily.xlsx, contains daily records of crude oil prices.
- **Time Period:** The dataset spans a significant period, capturing fluctuations in oil prices over time, helping to understand both short-term and long-term trends.
- **Key Variables:** It includes variables like date, daily closing price, opening price, highs, and lows.
- **Market Factors:** The dataset reflects the impact of various global events and economic indicators on crude oil prices.
- **Usage:** This data is used to analyze price trends, identify key drivers of price changes, and forecast future price movements.
- **Environmental and Economic Impact:** The data provides insights into the economic and environmental factors influencing crude oil prices.
- **products.** While they've seen success, we acknowledge the evolving market dynamics and the need to adapt our sales strategies to maintain and increase our competitive edge.



Exploratory Data Analysis



SUMMARY STATISTICS

```
[ ] df.describe()
```



	Date	Closing Value
count	8223	8216.000000
mean	2002-04-05 22:11:15.082086912	43.492139
min	1986-01-02 00:00:00	10.250000
25%	1994-01-25 12:00:00	19.577500
50%	2002-04-02 00:00:00	29.610000
75%	2010-06-12 12:00:00	63.402500
max	2018-07-09 00:00:00	145.310000
std	NaN	29.616804

Missing Values:

DATE	0
CLOSING VALUE	7


Replaced by Mean :

Number of Outliers: 34



```
df['Closing Value'].fillna(df['Closing Value'].mean(), inplace = True)
```


VISUALIZATIONS


Line plot



Histogram



Correlation HeatMap




Box Plot



Relational Plot

<Figure size 1000x600 with 0 Axes>





Model Building

We Implemented

ARIMA Model:

- Utilized for time-series analysis.
- Captures trends and seasonality effectively.

FBProphet Model:

- Handles missing data and outliers well.
- Suitable for business forecasting.

Linear Regression Model:

- Models relationships between variables.
- Simple and effective for predictive tasks.

LSTM (Long Short-Term Memory):

- Deep learning model for sequential data.
 - Captures long-term dependencies for accurate forecasting.
-

Arima

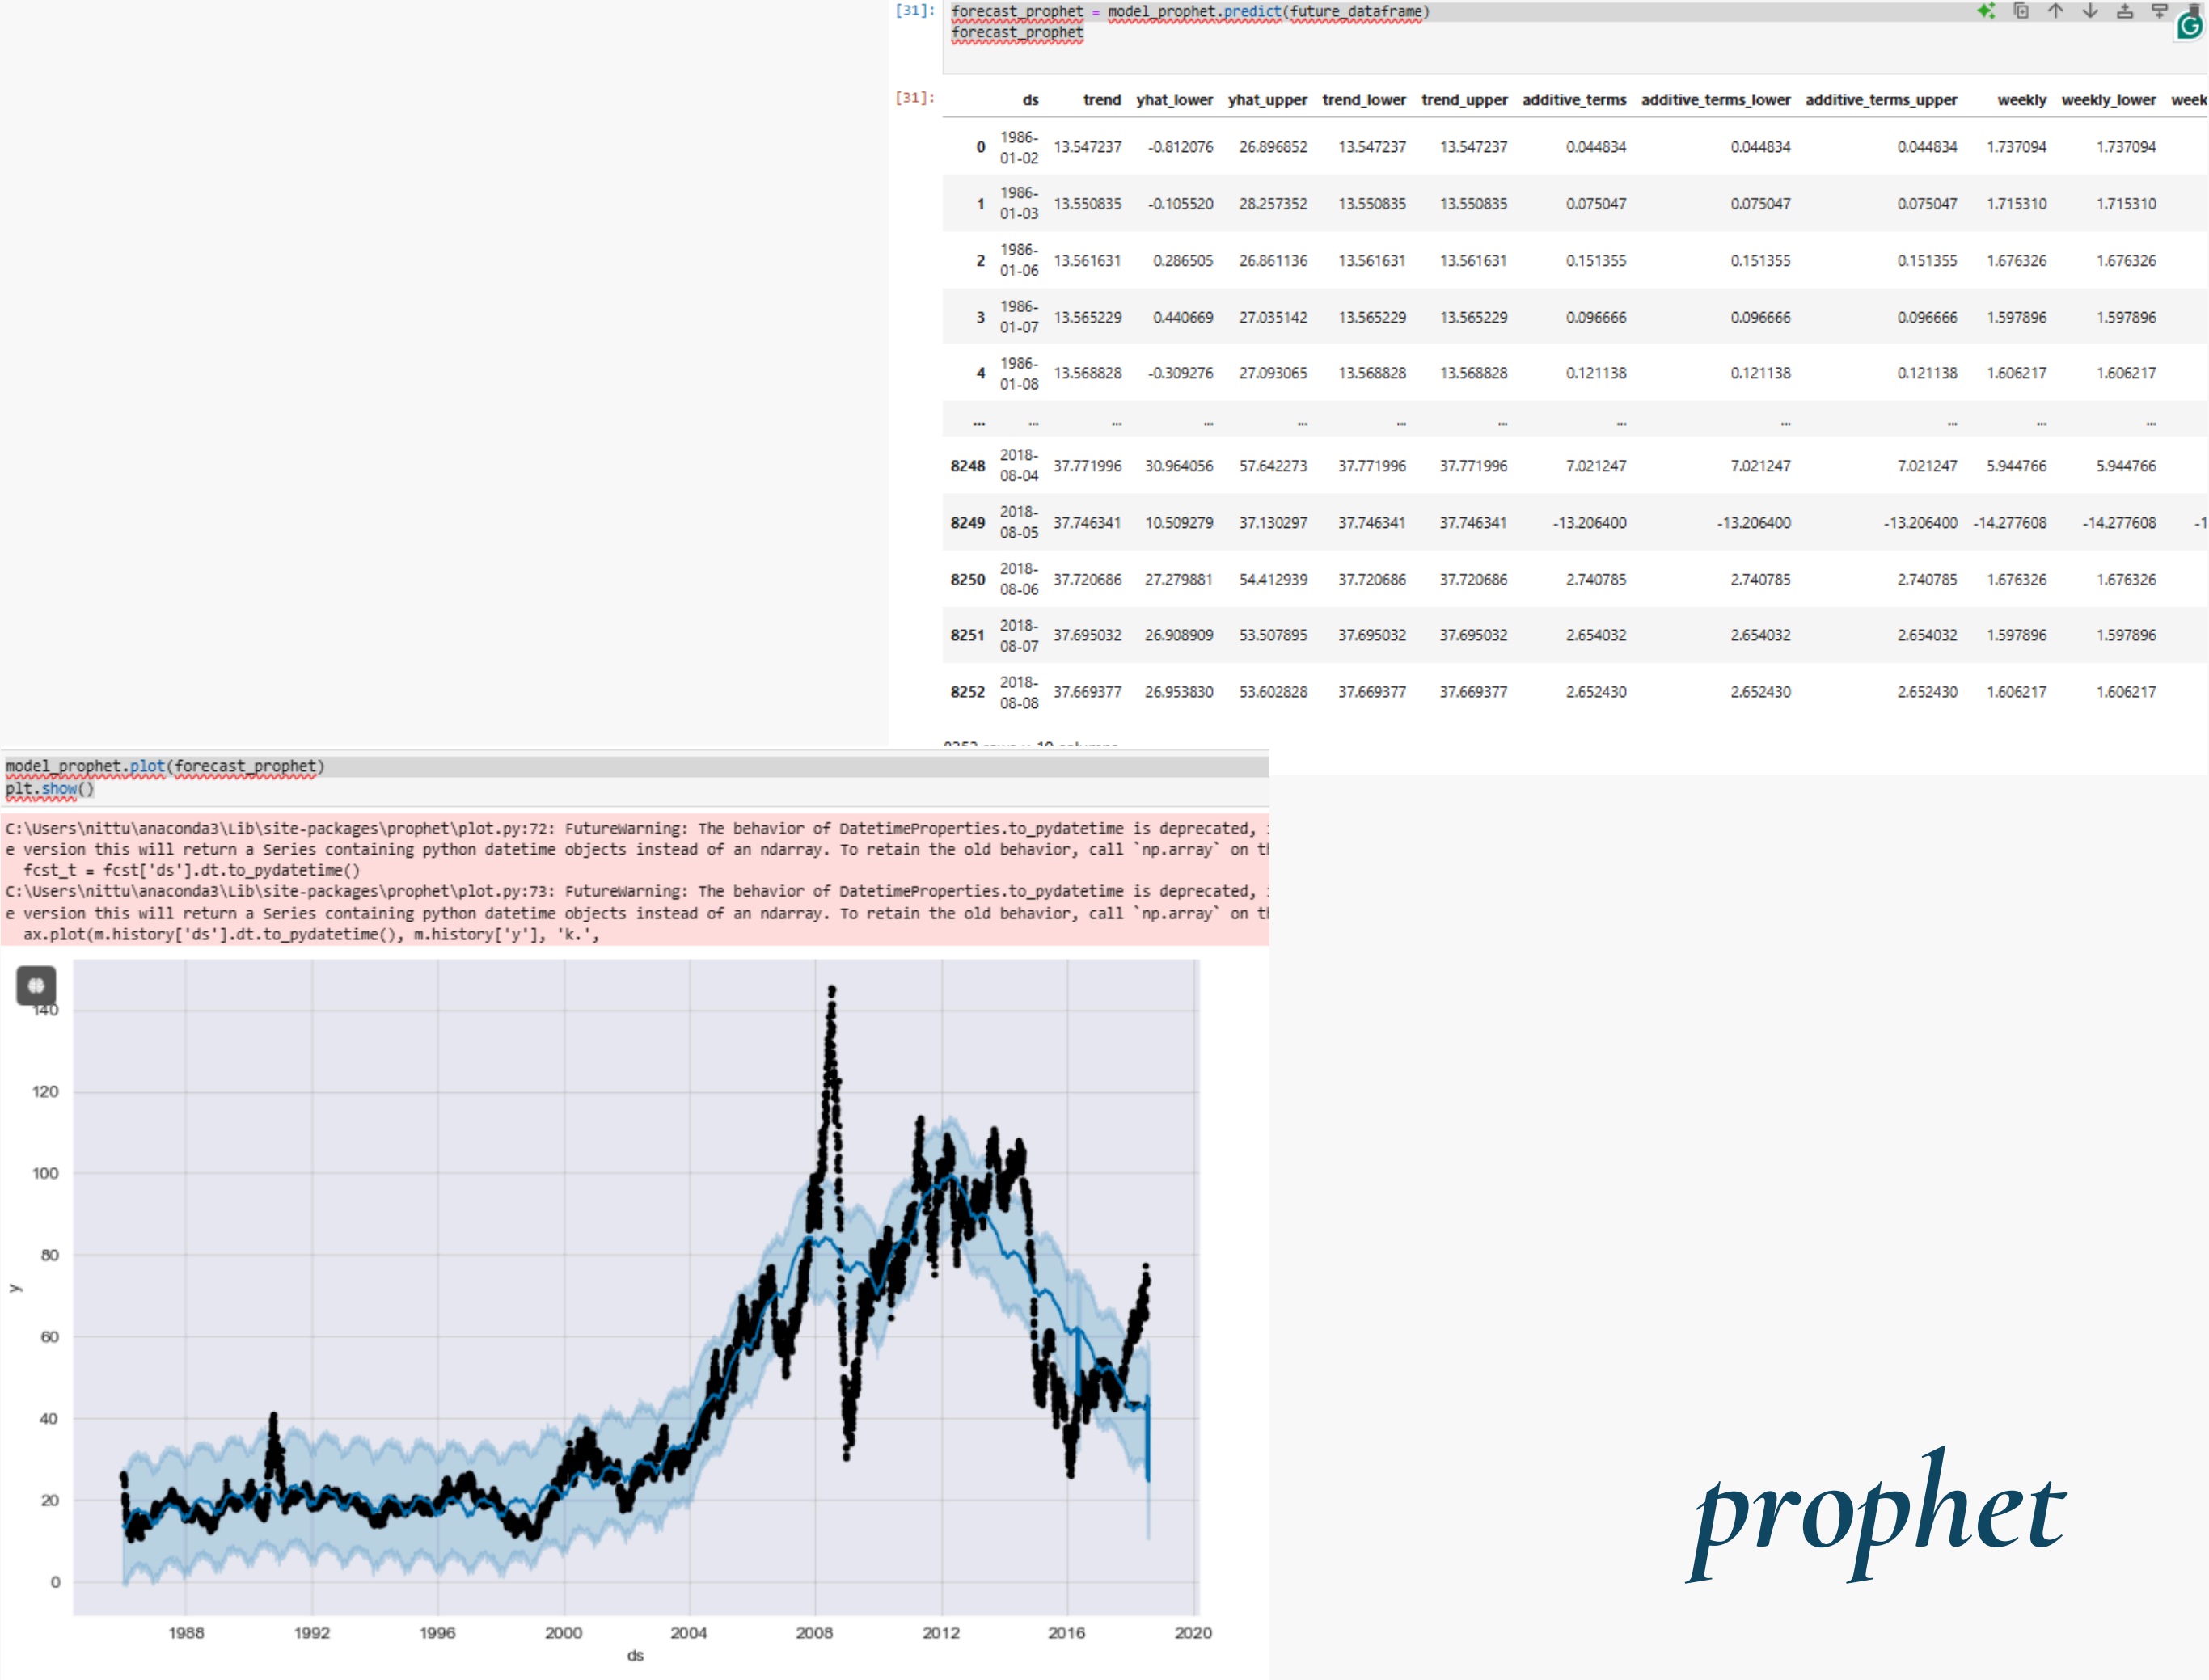
```
last_date = df['Date'].iloc[-1]
forecast_dates = pd.date_range(start=last_date + pd.DateOffset(days=1), periods=30)
forecast_df = pd.DataFrame({'Date': forecast_dates, 'Closing Value': forecast_arima})
print(forecast_df)
```

	Date	Closing Value
8223	2018-07-10	73.500150
8224	2018-07-11	73.777619
8225	2018-07-12	73.649127
8226	2018-07-13	73.662335
8227	2018-07-14	73.662736
8228	2018-07-15	73.665288
8229	2018-07-16	73.663272
8230	2018-07-17	73.664012
8231	2018-07-18	73.663908
8232	2018-07-19	73.663911
8233	2018-07-20	73.663895
8234	2018-07-21	73.663909
8235	2018-07-22	73.663904
8236	2018-07-23	73.663905
8237	2018-07-24	73.663905
8238	2018-07-25	73.663905
8239	2018-07-26	73.663905
8240	2018-07-27	73.663905
8241	2018-07-28	73.663905
8242	2018-07-29	73.663905
8243	2018-07-30	73.663905
8244	2018-07-31	73.663905
8245	2018-08-01	73.663905
8246	2018-08-02	73.663905
8247	2018-08-03	73.663905
8248	2018-08-04	73.663905
8249	2018-08-05	73.663905
8250	2018-08-06	73.663905
8251	2018-08-07	73.663905
8252	2018-08-08	73.663905

We implemented an ARIMA model to forecast the next 30 days of closing values. The model was fitted on historical data, and forecasted values were stored in a new DataFrame. And Plotted original closing values and ARIMA forecast for comparison.



Data Collection: We used the Prophet model for forecasting by first preparing the data with the required column names. After fitting the model, we created a future dataframe for the next 30 days and generated predictions. Finally, we plotted the forecast results to visualize the expected closing values.



prophet

Linear Regression

We implemented a Linear Regression model by converting the date into a numerical sequence representing days. After fitting the model using historical closing values, we predicted values for the next 30 days and visualized the results by plotting both the original data and the forecast.

```
plt.figure(figsize=(10, 6))
plt.plot(df['Days'], df['Closing Value'], label='Original')
plt.plot(future_days, forecast_lr, label='Linear Regression Forecast', color='green')
plt.legend()
plt.show()
```



LinearRegression

LinearRegression()

```
forecast_lr = model_lr.predict(future_days)
forecast_lr
```

C:\Users\nittu\anaconda3\Lib\site-packages\sklearn\base.py:493: UserWarning: Feature names

warnings.warn(

```
array([82.0481361 , 82.05462802, 82.06111994, 82.06761186, 82.074103
      82.0805957 , 82.08708761, 82.09357953, 82.10007145, 82.106563
      82.11305529, 82.11954721, 82.12603913, 82.13253105, 82.139022
      82.14551489, 82.15200681, 82.15849872, 82.16499064, 82.171482
      82.17797448, 82.1844664 , 82.19095832, 82.19745024, 82.203942
      82.21043408, 82.216926 , 82.22341792, 82.22990983, 82.236401
```

LSTM

We implemented an LSTM model for forecasting by first scaling the closing values using MinMaxScaler. We prepared the data by creating sequences of 60 time steps. The LSTM model was then built, trained, and used to predict future values. Finally, we visualized the forecasts from ARIMA, Prophet, Linear Regression, and LSTM models to compare their performances.




```
Epoch 1/10
256/256 ————— 25s 62ms/step - loss: 0.0110
Epoch 2/10
256/256 ————— 16s 62ms/step - loss: 0.0011
Epoch 3/10
256/256 ————— 17s 65ms/step - loss: 0.0011
Epoch 4/10
256/256 ————— 17s 67ms/step - loss: 0.0012
Epoch 5/10
256/256 ————— 16s 62ms/step - loss: 8.7179e-04
Epoch 6/10
256/256 ————— 16s 62ms/step - loss: 8.1532e-04
Epoch 7/10
256/256 ————— 16s 64ms/step - loss: 6.9748e-04
Epoch 8/10
256/256 ————— 16s 63ms/step - loss: 7.1644e-04
Epoch 9/10
256/256 ————— 16s 62ms/step - loss: 5.8990e-04
Epoch 10/10
256/256 ————— 15s 60ms/step - loss: 5.3796e-04
1/1 ————— 1s 778ms/step
```

Comparison between Models

This visualization presents a compact comparison of four forecasting models: ARIMA, Prophet, Linear Regression, and LSTM. Each model is evaluated using three key metrics: Mean Absolute Error (MAE), Mean Squared Error (MSE), and R² Score. The grouped bar chart allows for quick visual assessment of each model's performance, highlighting their strengths and weaknesses.

Model	MAE	MSE	R ² Score
ARIMA	6.13	67.67	-0.91
Prophet	28.64	862.75	-23.30
Linear Regression	14.15	235.64	-5.64
LSTM	14.15	235.64	-5.64





Deployment Using Streamlit App

ARIMA Model Selected For Deployment


Among all the models we have the lowest MAE and MSE with ARIMA Model and also the highest R2 score.

So here is the model_comparision with all models.

	Model	MAE	MSE	R2 Score
0	ARIMA	6.134121	67.673113	-0.906154
1	Prophet	28.639741	862.754576	-23.301275
2	Linear Regression	14.153864	235.638124	-5.637237
3	LSTM	14.153864	235.638124	-5.637237

After Deployment in Streamlit App

Crude Oil Price Forecasting App for ARIMA Model



Crude Oil

Select the start date for historical data:

1985/01/01

Select the end date for historical data:

2018/07/09

Enter the number of days to forecast:

30

Forecast


Enter the number of days to forecast:

30

Forecast

Forecasted Crude Oil Prices:

	Date	Forecasted Price
8,223	2018-07-10+00:00:00	73.8714
8,224	2018-07-11+00:00:00	73.9061
8,225	2018-07-12+00:00:00	73.9408
8,226	2018-07-13+00:00:00	73.9077
8,227	2018-07-14+00:00:00	73.9009
8,228	2018-07-15+00:00:00	73.9058
8,229	2018-07-16+00:00:00	73.9044
8,230	2018-07-17+00:00:00	73.902
8,231	2018-07-18+00:00:00	73.9036
8,232	2018-07-19+00:00:00	73.904



Explanation

In the Streamlit deployment, the ARIMA model is used to forecast crude oil prices for a user-specified number of days.

The user inputs the number of days they want to forecast, and upon clicking the "Forecast" button, the model makes predictions using the `predict_crude_oil_price` function.

The results, both the numerical values and a visual plot of the forecast, are then displayed in the Streamlit app.

The model's deployment in Streamlit allows for easy access and use of the forecasting functionality.

Users can simply input the number of days they are interested in and obtain the predicted prices without needing to interact with any complex code.

The visual plot further enhances the ease of understanding and interpretation of the forecasts.

Thank You